

NC STATE UNIVERSITY

Programming in R Part II

Justin Post

August 9, 2017

What do we want to be able to do?

- Restructure Data/Clean Data
- Streamline repeated sections of code
- **Improve efficiency of code**
- Write custom functions to simplify code

Efficient Code

For loops inefficient in R

- R interpreted language
- Must figure out how to evaluate code at each iteration of loop
- Slows it down

Efficient Code

For loops inefficient in R

- R interpreted language
- Must figure out how to evaluate code at each iteration of loop
- Slows it down

Vecotrized functions much faster!

- Vectorized function: works on entire vector at once
- Avoids costly computation time

Efficient Code

Some 'built-in' vectorized functions

- `colMeans()`, `rowMeans()`
- `colSums()`, `rowSums()`
- `colSds()`, `colVars()`, `colMedians()` (`matrixStats` package)
- `ifelse()`
- `apply()` family
- Create your own with `Vectorize()`

Efficient Code

- Find column means for full Batting data set
- `colMeans()` just requires a numeric data frame (array)

```
colMeans(select(Batting, G:GIDP), na.rm = TRUE)
```

##	G	AB	R	H	X2B	X3B
##	51.400111	149.970327	19.887038	39.261647	6.637067	1.373361
##	HR	RBI	SB	CS	BB	SO
##	2.949305	17.965163	3.158184	1.324025	13.811484	21.629849
##	IBB	HBP	SH	SF	GIDP	
##	1.213234	1.113395	2.457900	1.150122	3.210032	

Efficient Code

- Compare computational time
- `microbenchmark` package allows for easy recording of computing time

```
install.packages("microbenchmark")
```

```
library(microbenchmark)
```

Efficient Code

- Compare computational time

```
microbenchmark(
  colMeans(select(Batting, G:GIDP), na.rm = TRUE)
)
```

```
## Unit: milliseconds
```

```
##              expr      min       lq    mean
## colMeans(select(Batting, G:GIDP), na.rm = TRUE) 10.5278 11.00802 13.866
##      median      uq      max neval
## 11.46476 15.06337 67.50874   100
```


Efficient Code

- Compare computational time

```
microbenchmark(
  for(i in 1:17){
    mean(Batting[, i + 5][[1]], na.rm = TRUE)
  }
)
```

```
## Unit: milliseconds
```

```
##                                     expr      min
## for (i in 1:17) { mean(Batting[, i + 5][[1]], na.rm = TRUE) } 2.67567
##      lq      mean    median      uq      max neval
## 2.909959 3.303627 3.092535 3.423539 6.953662   100
```

Efficient Code

- With vectorized functions, can easily find cool stuff
 - Median number of games played for all players
 - Median number of AB for players that batted
 - Steps: (think `dplyr` commands!)
1. Group observations by `playerID`
 2. Summarise variables of interest
 3. Remove non numeric column
 4. Coerce to matrix for use in `colMedians()`
 5. Use `colMedians()` function

Efficient Code

```
library(matrixStats) #install if not installed

##
## Attaching package: 'matrixStats'

## The following object is masked from 'package:dplyr':
##
##      count

Batting %>% group_by(playerID) %>%
  summarise(totG = sum(G), totAB = sum(AB)) %>%
  select(-playerID) %>% as.matrix() %>%
  colMedians(na.rm = TRUE)

## [1] 78 91
```

Efficient Code

- Logical statement - comparison of two quantities
 - resolves as TRUE or FALSE
- Often want to execute code logically
- logical comparison operators
 - `==`, `!=`, `>=`, `<=`, `>`, `<`
 - `&` "and"
 - `|` "or"
- logical functions
 - `is.` family (`is.numeric()`, `is.data.frame()`, etc.)

Efficient Code

If then, If then else

- Often want to execute statements conditionally
- If then else concept

```
if (condition) {  
  then execute code  
}
```

```
#if then else  
if (condition) {  
  execute this code  
} else {  
  execute this code  
}
```

Efficient Code

If then, If then else

- Often want to execute statements conditionally
- If then else concept

#Or more if statements

```
if (condition) {  
  execute this code  
} else if (condition2) {  
  execute this code  
} else if (condition3) {  
  execute this code  
} else {  
  #if no conditions met  
  execute this code  
}
```

Efficient Code

- Often create new variables

Efficient Code

- Often create new variables
- Built in data set `airquality`
 - daily air quality measurements in New York
 - from May (Day 1) to September (Day 153) in 1973

```
airquality<-tbl_df(airquality)
airquality
```

```
## # A tibble: 153 x 6
##   Ozone Solar.R  Wind  Temp Month   Day
##   <int>   <int> <dbl> <int> <int> <int>
## 1    41    190   7.4    67     5     1
## 2    36    118   8.0    72     5     2
## 3    12    149  12.6    74     5     3
## 4    18    313  11.5    62     5     4
## 5     NA     NA  14.3    56     5     5
## # ... with 148 more rows
```


Efficient Code

Want to code a wind category variable

- high wind days ($15\text{mph} \leq \text{wind}$)
- windy days ($10\text{mph} \leq \text{wind} < 15\text{mph}$)
- lightwind days ($6\text{mph} \leq \text{wind} < 10\text{mph}$)
- calm days ($\text{wind} \leq 6\text{mph}$)

Efficient Code

Want to code a wind category variable

- high wind days ($15\text{mph} \leq \text{wind}$)
- windy days ($10\text{mph} \leq \text{wind} < 15\text{mph}$)
- lightwind days ($6\text{mph} \leq \text{wind} < 10\text{mph}$)
- calm days ($\text{wind} \leq 6\text{mph}$)

Initial plan

- loop through each observation
- use if then else to determine wind status

Efficient Code

#initialize vector to save results

```
status<-vector()
```

```
for (i in 1:(dim(airquality)[1])){  
  if(airquality$Wind[i] >= 15){  
    status[i] <- "HighWind"  
  } else if (airquality$Wind[i] >= 10){  
    status[i] <- "Windy"  
  } else if (airquality$Wind[i] >= 6){  
    status[i] <- "LightWind"  
  } else if (airquality$Wind[i] >= 0){  
    status[i] <- "Calm"  
  } else {  
    status[i] <- "Error"  
  }  
}
```

Efficient Code

status

```
## [1] "LightWind" "LightWind" "Windy"      "Windy"      "Windy"
## [6] "Windy"      "LightWind" "Windy"      "HighWind"   "LightWind"
## [11] "LightWind" "LightWind" "LightWind"  "Windy"      "Windy"
## [16] "Windy"      "Windy"      "HighWind"   "Windy"      "LightWind"
## [21] "LightWind" "HighWind"   "LightWind"  "Windy"      "HighWind"
## [26] "Windy"      "LightWind" "Windy"      "Windy"      "Calm"
## [31] "LightWind" "LightWind" "LightWind"  "HighWind"   "LightWind"
## [36] "LightWind" "Windy"      "LightWind"  "LightWind"  "Windy"
## [41] "Windy"      "Windy"      "LightWind"  "LightWind"  "Windy"
## [46] "Windy"      "Windy"      "HighWind"   "LightWind"  "Windy"
## [51] "Windy"      "LightWind" "Calm"        "Calm"        "LightWind"
## [56] "LightWind" "LightWind" "Windy"      "Windy"      "Windy"
## [61] "LightWind" "Calm"       "LightWind"  "LightWind"  "Windy"
## [66] "Calm"       "Windy"      "Calm"        "LightWind"  "Calm"
## [71] "LightWind" "LightWind" "Windy"      "Windy"      "Windy"
## [76] "Windy"      "LightWind" "Windy"      "LightWind"  "Calm"
## [81] "Windy"      "LightWind" "LightWind"  "Windy"      "LightWind"
## [86] "LightWind" "LightWind" "Windy"      "LightWind"  "LightWind"
```

20/84

Efficient Code

- Add it to the data set

```
airquality$status <- status
```

- Find mean temperature for each wind Status

```
airquality$status <- status  
airquality %>% group_by(status) %>%  
  mutate(avgTemp = mean(Temp))
```

Efficient Code

```
## # A tibble: 153 x 8
## # Groups:   status [4]
##   Ozone Solar.R Wind Temp Month Day status avgTemp
##   <int> <int> <dbl> <int> <int> <int> <chr> <dbl>
## 1 41 190 7.4 67 5 1 LightWind 79.43077
## 2 36 118 8.0 72 5 2 LightWind 79.43077
## 3 12 149 12.6 74 5 3 Windy 75.54839
## 4 18 313 11.5 62 5 4 Windy 75.54839
## 5 NA NA 14.3 56 5 5 Windy 75.54839
## # ... with 148 more rows
```

Efficient Code

- Know for loops not great
- `ifelse()` is vectorized version of `if then else`
- Syntax

```
ifelse(vector_condition, if_true_do_this, if_false_do_this)
```

Efficient Code

```
ifelse(airquality$Wind >= 15, "HighWind",
       ifelse(airquality$Wind >= 10, "Windy",
              ifelse(airquality$Wind >= 6, "LightWind", "Calm"))))
```

```
## [1] "LightWind" "LightWind" "Windy"      "Windy"      "Windy"
## [6] "Windy"      "LightWind" "Windy"      "HighWind"    "LightWind"
## [11] "LightWind" "LightWind" "LightWind"  "Windy"      "Windy"
## [16] "Windy"      "Windy"      "HighWind"   "Windy"      "LightWind"
## [21] "LightWind" "HighWind"   "LightWind"  "Windy"      "HighWind"
## [26] "Windy"      "LightWind" "Windy"      "Windy"      "Calm"
## [31] "LightWind" "LightWind" "LightWind"  "HighWind"    "LightWind"
## [36] "LightWind" "Windy"      "LightWind"  "LightWind"   "Windy"
## [41] "Windy"      "Windy"      "LightWind"  "LightWind"   "Windy"
## [46] "Windy"      "Windy"      "HighWind"   "LightWind"   "Windy"
## [51] "Windy"      "LightWind"  "Calm"       "Calm"        "LightWind"
## [56] "LightWind" "LightWind" "Windy"      "Windy"       "Windy"
## [61] "LightWind" "Calm"       "LightWind"  "LightWind"   "Windy"
## [66] "Calm"       "Windy"      "Calm"       "LightWind"   "Calm"
## [71] "LightWind" "LightWind" "Windy"      "Windy"       "Windy"
## [76] "Windy"      "LightWind" "Windy"      "LightWind"   "Calm"
```

24/84

Efficient Code

- Compare speed

```
loopTime<-microbenchmark(  
  for (i in 1:(dim(airquality)[1])){  
    if(airquality$Wind[i] >= 15){  
      status[i] <- "HighWind"  
    } else if (airquality$Wind[i] >= 10){  
      status[i] <- "Windy"  
    } else if (airquality$Wind[i] >= 6){  
      status[i] <- "LightWind"  
    } else if (airquality$Wind[i] >= 0){  
      status[i] <- "Calm"  
    } else{  
      status[i] <- "Error"  
    }  
  }  
, unit = "us")
```

Efficient Code

- Compare speed

```
vectorTime <- microbenchmark(  
  ifelse(airquality$Wind >= 15, "HighWind",  
    ifelse(airquality$Wind >= 10, "Windy",  
      ifelse(airquality$Wind >= 6, "LightWind", "Calm"))))  
, unit = "us")
```

Efficient Code (Note units!)

loopTime

```
## Unit: microseconds
##
## for (i in 1:(dim(airquality)[1])) {      if (airquality$Wind[i] >= 15) {      status[i]
##      min      lq      mean  median      uq      max neval
## 29039.58 31981.91 34859.96 33797.79 36785.91 61407.75 100
```

vectorTime

```
## Unit: microseconds
##
## ifelse(airquality$Wind >= 15, "HighWind", ifelse(airquality$Wind >= 10, "Windy", if
##      min      lq      mean  median      uq      max neval
## 287.385 400.0875 691.6916 555.6225 670.6945 6716.412 100
```

Efficient Code

- `apply()` family of functions *pretty* fast
- Check `help(apply)`
 - We'll look at `apply()`, `sapply()`, `lapply()`

Efficient Code

- `apply()` family of functions *pretty* fast
- Check `help(apply)`
 - We'll look at `apply()`, `sapply()`, `lapply()`
 - Use `apply()` to find summary for columns of `airquality` data

```
apply(X = select(airquality, Ozone:Temp), MARGIN = 2,  
      FUN = summary, na.rm = TRUE)
```

Efficient Code

- Keeps data numeric, keeps labels!

```
## $Ozone
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##      1.00   18.00   31.50   42.13   63.25   168.00    37
##
## $Solar.R
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##       7.0   115.8   205.0   185.9   258.8   334.0     7
##
## $Wind
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     1.700   7.400   9.700   9.958  11.500   20.700
##
## $Temp
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     56.00   72.00   79.00   77.88   85.00   97.00
```

Efficient Code

- Use `lapply()` to apply function to lists
- Obtain a list object

```
fit <- lm(Ozone ~ Wind, data = airquality)
fit <- list(fit$residuals, fit$effects, fit$fitted.values)
```

Efficient Code

```
fit[[1]]
```

```
##           1           2           3           4           6           7
## -14.7960653 -16.4655116 -14.9312663 -15.0372815  13.8358563 -26.1349578
##           8           9          11          12          13          14
##  -1.2701589  22.7006553 -51.5715267 -27.0289427 -34.8044041 -22.3678352
##           15          16          17          18          19          20
##  -5.6007126 -19.0372815   3.7381799  11.2640864  -3.0372815 -32.0289427
##          21          22          23          24          28          29
## -42.0289427   6.2724252 -39.0289427   1.7381799  -7.2618201  30.8358563
##          30          31          38          40          41          44
##  49.7673658 -18.7960653 -14.0289427  50.7298411   5.9627185 -29.4655116
##          47          48          49          50          51          62
##   6.8358563  55.0312090 -25.8044041 -21.0372815 -26.6983889  60.8858892
##          63          64          66          67          68          69
##   3.1955959 -13.8044041  -7.3386494   3.6321648   8.4368121  35.0979195
##          70          71          73          74          76          77
##  31.7673658  29.2039347  -7.4946974  12.8358563 -10.4946974 -10.5715267
##          78          79          80          81          82          85
##  -4.6983889  -0.9020805  10.4368121  29.9627185 -42.5715267  30.8650422
```

32/84

Efficient Code

```
fit[[2]]
```

```
## (Intercept)      Wind
## -453.7465588 -212.8004841 -14.7950343 -14.4827016 13.0973609
##
## -24.4774608 -1.5903064 19.9845153 -49.2674921 -25.7897935
##
## -33.3750969 -21.5850656 -5.6926704 -18.4827016 4.1026017
##
## 9.1944841 -2.4827016 -30.7897935 -40.7897935 4.8873922
##
## -37.7897935 2.1026017 -6.8973983 30.0973609 52.5277800
##
## -16.6821888 -12.7897935 50.4096936 6.5172984 -27.5798248
##
## 6.0973609 52.0868793 -24.3750969 -20.4827016 -25.6874296
##
## 64.2548094 4.6249031 -12.3750969 -4.1598873 4.4149344
##
## 11.4254160 37.6301439 34.5277800 31.3178112 -8.0050031
```

33/84

Efficient Code

- Apply `mean()` function to each list element

```
lapply(X = fit, FUN = mean)
```

```
## [[1]]  
## [1] -5.731915e-16  
##  
## [[2]]  
## [1] -4.333566  
##  
## [[3]]  
## [1] 42.12931
```

Efficient Code

- Use `sapply()` similar but returns a vector if possible

```
sapply(X = fit, FUN = mean)
```

```
## [1] -5.731915e-16 -4.333566e+00 4.212931e+01
```

- `apply()` functions not as good as `colMeans()` type functions

```
air2 <- select(airquality, Ozone:Day)
microbenchmark(apply(X = air2, MARGIN = 2, FUN = mean, na.rm = TRUE))
```

```
## Unit: microseconds
```

```
##                               expr      min       lq
##  apply(X = air2, MARGIN = 2, FUN = mean, na.rm = TRUE) 135.008 138.9555
##      mean   median      uq      max neval
## 193.8308 142.7055 161.4565 1371.389   100
```

```
microbenchmark(colMeans(air2, na.rm = TRUE))
```

```
## Unit: microseconds
```

```
##                               expr      min       lq      mean median      uq      max
##  colMeans(air2, na.rm = TRUE)  58.425  60.004  64.19629  60.794  62.57 187.511
##  neval
##      100
```

Recap!

- Vectorized functions fast!
- 'Built-in' vectorized functions
 - `colMeans()`, `rowMeans()`
 - `colSums()`, `rowSums()`
 - `colSds()`, `colVars()`, `colMedians()` (`matrixStats` package)
 - `ifelse()`
 - `apply()` family

Activity

- [**Vectorized Functions Activity** instructions](#) available on web
- Work in small groups
- Ask questions! TAs and I will float about the room
- Feel free to ask questions about anything you didn't understand as well!

What do we want to be able to do?

- Restructure Data/Clean Data
- Streamline repeated sections of code
- Improve efficiency of code
- **Write custom functions to simplify code**

Writing Functions

- Knowing how to write **functions** vital to custom analyses!
- Function writing syntax

```
nameOfFunction <- function(input1, input2, ...) {  
  #code  
  #return something with return()  
  #or returns last value  
}
```


Writing Functions

- Can look at code for functions

var

```
## function (x, y = NULL, na.rm = FALSE, use)
## {
##   if (missing(use))
##     use <- if (na.rm)
##       "na.or.complete"
##     else "everything"
##   na.method <- pmatch(use, c("all.obs", "complete.obs", "pairwise.complete.obs",
##     "everything", "na.or.complete"))
##   if (is.na(na.method))
##     stop("invalid 'use' argument")
##   if (is.data.frame(x))
##     x <- as.matrix(x)
##   else stopifnot(is.atomic(x))
##   if (is.data.frame(y))
##     y <- as.matrix(y)
##   else stopifnot(is.atomic(y))
```

41/84

Writing Functions

- Can look at code for functions

colMeans

```
## function (x, na.rm = FALSE, dims = 1L)
## {
##   if (is.data.frame(x))
##     x <- as.matrix(x)
##   if (!is.array(x) || length(dn <- dim(x)) < 2L)
##     stop("'x' must be an array of at least two dimensions")
##   if (dims < 1L || dims > length(dn) - 1L)
##     stop("invalid 'dims'")
##   n <- prod(dn[id <- seq_len(dims)])
##   dn <- dn[-id]
##   z <- if (is.complex(x))
##     .Internal(colMeans(Re(x), n, prod(dn), na.rm)) + (0+1i) *
##     .Internal(colMeans(Im(x), n, prod(dn), na.rm))
##   else .Internal(colMeans(x, n, prod(dn), na.rm))
##   if (length(dn) > 1L) {
##     dim(z) <- dn
```

42/84

Writing Functions

- Can look at code for functions

mean

```
## function (x, ...)  
## UseMethod("mean")  
## <bytecode: 0x0000000018519a90>  
## <environment: namespace:base>
```

Writing Functions

- Can look at code for functions

mean.default

```
## function (x, trim = 0, na.rm = FALSE, ...)  
## {  
##     if (!is.numeric(x) && !is.complex(x) && !is.logical(x)) {  
##         warning("argument is not numeric or logical: returning NA")  
##         return(NA_real_)  
##     }  
##     if (na.rm)  
##         x <- x[!is.na(x)]  
##     if (!is.numeric(trim) || length(trim) != 1L)  
##         stop("'trim' must be numeric of length one")  
##     n <- length(x)  
##     if (trim > 0 && n) {  
##         if (is.complex(x))  
##             stop("trimmed means are not defined for complex data")  
##         if (anyNA(x))  
##             return(NA_real_)
```

44/84

Writing Functions

- Goal: Create a `standardize()` function
- Take vector of values
 - subtract mean
 - divide by standard deviation
- z-score idea
- Formula: For value i ,
$$(\text{value}[i] - \text{mean}(\text{value})) / \text{sd}(\text{value})$$

Writing Functions

```
nameOfFunction <- function(input1, input2, ...) {  
  #code  
  #return something with return()  
  #or returns last value  
}
```

```
standardize <- function(vector) {  
  return((vector - mean(vector)) / sd(vector))  
}
```

Writing Functions

- Now use it!

```
data <- runif(5)
```

```
data
```

```
## [1] 0.40420982 0.16632350 0.82994597 0.27298712 0.02134729
```

```
result <- standardize(data)
```

```
result
```

```
## [1] 0.2115897 -0.5598517 1.5922093 -0.2139523 -1.0299950
```

Writing Functions

- Check result has mean 0 and sd 1

```
mean(result)
```

```
## [1] -7.769393e-17
```

```
sd(result)
```

```
## [1] 1
```


Writing Functions

- Goal: Add more inputs
- Make centering optional
- Make scaling optional

```
standardize <- function(vector, center, scale) {  
  if (center == TRUE) {  
    vector <- vector - mean(vector)  
  }  
  if (scale == TRUE) {  
    vector <- vector / sd(vector)  
  }  
  return(vector)  
}
```

Writing Functions

```
result <- standardize(data, center = TRUE, scale = TRUE)  
result
```

```
## [1] 0.2115897 -0.5598517 1.5922093 -0.2139523 -1.0299950
```

```
result <- standardize(data, center = FALSE, scale = TRUE)  
result
```

```
## [1] 1.31081187 0.53937041 2.69143144 0.88526981 0.06922711
```

Writing Functions

- Give center and scale default arguments

```
standardize <- function(vector, center = TRUE, scale = TRUE) {  
  #center and scale if appropriate  
  if (center == TRUE) {  
    vector <- vector - mean(vector)  
  }  
  if (scale == TRUE) {  
    vector <- vector / sd(vector)  
  }  
  return(vector)  
}
```

Writing Functions

```
result <- standardize(data, center = TRUE, scale = TRUE)
result
```

```
## [1]  0.2115897 -0.5598517  1.5922093 -0.2139523 -1.0299950
```

#same call

```
result <- standardize(data)
result
```

```
## [1]  0.2115897 -0.5598517  1.5922093 -0.2139523 -1.0299950
```

Writing Functions

- Return more than 1 object by returning a list
- Goal: Also return
 - `mean()` of original data
 - `sd()` of original data

Writing Functions

```
standardize <- function(vector, center = TRUE, scale = TRUE) {  
  #get attributes to return  
  mean <- mean(vector)  
  stdev <- sd(vector)  
  #center and scale if appropriate  
  if (center == TRUE) {  
    vector <- vector - mean  
  }  
  if (scale == TRUE) {  
    vector <- vector / stdev  
  }  
  #return a list of objects  
  return(list(vector, mean, stdev))  
}
```

Writing Function

```
result <- standardize(data)
result
```

```
## [[1]]
## [1]  0.2115897 -0.5598517  1.5922093 -0.2139523 -1.0299950
##
## [[2]]
## [1] 0.3389627
##
## [[3]]
## [1] 0.308366
```

```
result[[2]]
```

```
## [1] 0.3389627
```

Writing Functions

- Fancy up what we return by giving names

```
standardize <- function(vector, center = TRUE, scale = TRUE) {  
  #get attributes to return  
  mean <- mean(vector)  
  stdev <- sd(vector)  
  #center and scale if appropriate  
  if (center == TRUE) {  
    vector <- vector - mean  
  }  
  if (scale == TRUE) {  
    vector <- vector / stdev  
  }  
  #return a list of objects  
  return(list(result = vector, mean = mean, sd = stdev))  
}
```


Writing Functions

```
result <- standardize(data, center = TRUE, scale = TRUE)
result
```

```
## $result
## [1]  0.2115897 -0.5598517  1.5922093 -0.2139523 -1.0299950
##
## $mean
## [1] 0.3389627
##
## $sd
## [1] 0.308366
```

```
result$sd
```

```
## [1] 0.308366
```

Writing Functions

- Can bring in unnamed arguments
- Arguments that can be used by functions **inside** your function
- Done already in `apply()`

`apply`

```
## function (X, MARGIN, FUN, ...)  
## {  
##     FUN <- match.fun(FUN)  
##     dl <- length(dim(X))  
##     if (!dl)  
##         stop("dim(X) must have a positive length")  
##     if (is.object(X))  
##         X <- if (dl == 2L)  
##             as.matrix(X)  
##             else as.array(X)  
##     d <- dim(X)  
##     dn <- dimnames(X)  
##     ds <- seq_len(dl)
```

58/84

Writing Functions

```
apply(X = select(airquality, Ozone:Temp), MARGIN = 2,
      FUN = summary, na.rm = TRUE)
```

```
## $Ozone
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##   1.00  18.00   31.50   42.13  63.25  168.00    37
##
## $Solar.R
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##    7.0  115.8   205.0   185.9  258.8   334.0     7
##
## $Wind
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  1.700   7.400   9.700   9.958  11.500   20.700
##
## $Temp
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  56.00   72.00   79.00   77.88   85.00   97.00
```

Writing Functions

- Add unnamed arguments to our function

```
standardize <- function(vector, center = TRUE, scale = TRUE, ...) {  
  #get attributes to return  
  mean <- mean(vector, ...)  
  stdev <- sd(vector, ...)  
  #center and scale if appropriate  
  if (center == TRUE) {  
    vector <- vector - mean  
  }  
  if (scale == TRUE) {  
    vector <- vector / stdev  
  }  
  #return a list of objects  
  return(list(result = vector, mean = mean, sd = stdev))  
}
```

Writing Functions

```
sData <- standardize(airquality$Ozone, na.rm = TRUE)
```

```
sData$mean
```

```
## [1] 42.12931
```

```
sData$sd
```

```
## [1] 32.98788
```

```
sData$result
```

```
## [1] -0.03423409 -0.18580489 -0.91334473 -0.73145977 NA
## [6] -0.42831817 -0.57988897 -0.70114561 -1.03460136 NA
## [11] -1.06491552 -0.79208809 -0.94365889 -0.85271641 -0.73145977
## [16] -0.85271641 -0.24643321 -1.09522968 -0.36768985 -0.94365889
## [21] -1.24680048 -0.94365889 -1.15585800 -0.30706153 NA
## [26] NA NA -0.57988897 0.08702254 2.20901373
## [31] -0.15549073 NA NA NA NA NA
```

61/84

Recap!

- Function writing opens R up!
- Syntax

```
nameOfFunction <- function(input1, input2, ...) {  
  #code  
  #return something with return()  
  #or returns last value  
}
```

- Can set defaults in function definition
- Can return a named list
- Can give unnamed arguments for use

Activity

- [Function Writing Activity instructions](#) available on web
- Work in small groups
- Ask questions! TAs and I will float about the room
- Feel free to ask questions about anything you didn't understand as well!

What do we want to be able to do?

- Restructure Data/Clean Data
- Streamline repeated sections of code
- **Improve efficiency of code**
- Write custom functions to simplify code

Parallel Computing

- Just basic intro and idea (very complicated, many things to consider!)

Idea:

- Take computations that can be done independently (or close to it)
- Don't run sequentially
- Split up computation
- Run computation simultaneously on
 - different processor cores
 - across many connected computers (i.e. on a cluster)
 - or a few other ways
- Combine results

Parallel Computing

Many applications in data science lend themselves to parallel computing

Examples

- Monte Carlo simulation studies
- Bootstrapping
- Multiple MCMC runs from different starting points
- Cross Validation
- Random Forests and Boosting algorithms
- We'll use `parallel` package (built-in)

Parallel Computing

- `parallel` package function we'll use has syntax similar to `apply()` family
- Problem to parallelize:
 - kmeans clustering
 - group similar observations
 - consider iris data set

```
iris<-tbl_df(iris)
```

Parallel Computing

iris


```
## # A tibble: 150 x 5
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
##         <dbl>         <dbl>         <dbl>         <dbl>   <fctr>
## 1         5.1           3.5           1.4           0.2   setosa
## 2         4.9           3.0           1.4           0.2   setosa
## 3         4.7           3.2           1.3           0.2   setosa
## 4         4.6           3.1           1.5           0.2   setosa
## 5         5.0           3.6           1.4           0.2   setosa
## # ... with 145 more rows
```

Parallel Computing

- Problem to parallelize:
 - kmeans clustering
 - group similar observations

```
library(fpc)          #install if needed  
iris$Species <- NULL #remove category labels (truly 3 groups)  
clus <- kmeans(iris, centers = 3, nstart = 100)  
plotcluster(iris, clus$cluster)
```

Parallel Computing



Parallel Computing

Why is this able to be parallelized?

- Code tries to find 3 'cluster' centers using 100 random starting positions
- Result is the starting position that yields the minimal `result$tot.withinss` value
- Each random starting point used is independent of the others (embarrassingly parallel)
- Can assign some of the runs of the algorithm to separate computer cores
- Combine back at the end, look at overall smallest value

Parallel Computing

- How to parallelize this?
- Create a function using lapply to do the kmeans call

```
parallel.function <- function(data, i) {  
  kmeans(as.matrix(data), centers = 3, nstart = i)  
}
```


Parallel Computing

- Evaluating example with lapply

```
results <- lapply(X = c(25, 25), FUN = parallel.function, data = iris)
```

Parallel Computing

```
results[[1]]
```

```
## K-means clustering with 3 clusters of sizes 38, 50, 62
```

##

Cluster means:

```
## Sepal.Length Sepal.Width Petal.Length Petal.Width
```

```
## 1      6.850000      3.073684      5.742105      2.071053
```

##	2	5.006000	3.428000	1.462000	0.246000
----	---	----------	----------	----------	----------

```
## 3      5.901613      2.748387      4.393548      1.433871
```

##

```
## Clustering vector:
```

```
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
## [36] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 3 3 1 3 3 3 3 3 3 3 3 3 3 3 3
```

```
## [71] 3 3 3 3 3 3 3 1 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 1 3 1 1 1
```

[106] 1 3 1 1 1 1 1 1 3 3 1 1 1 1 3 1 3 1 3 1 1 3 3 1 1 1 1 1 3 1 1 1 1 3 1

```
## [141] 1 1 3 1 1 1 3 1 1 3
```

##

```
## Within cluster sum of squares by cluster:
```

```
## [1] 23.87947 15.15100 39.82097
```

```
## (between SS / total SS = 88.4 %)
```

Parallel Computing

```
results[[2]]
```

```
## K-means clustering with 3 clusters of sizes 38, 50, 62
```

```
##
```

```
## Cluster means:
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width
```

```
## 1      6.850000      3.073684      5.742105      2.071053
```

```
## 2      5.006000      3.428000      1.462000      0.246000
```

```
## 3      5.901613      2.748387      4.393548      1.433871
```

```
##
```

```
## Clustering vector:
```

```
##   [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
##  [36] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 3 3 1 3 3 3 3 3 3 3 3 3 3 3 3 3 3
```

```
##  [71] 3 3 3 3 3 3 1 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 1 3 1 1 1
```

```
## [106] 1 3 1 1 1 1 1 1 3 3 1 1 1 1 3 1 3 1 3 1 1 3 3 1 1 1 1 1 3 1 1 1 3 1
```

```
## [141] 1 1 3 1 1 1 3 1 1 3
```

```
##
```

```
## Within cluster sum of squares by cluster:
```

```
## [1] 23.87947 15.15100 39.82097
```

```
## (between_SS / total_SS =  88.4 %)
```

75/84

Parallel Computing

- Want best result of the two returned

Parallel Computing

- Want best result of the two returned
- Create a function to determine which call had the best overall result

Parallel Computing

- Set-up is a lot of work!
- For a large problem this could save a lot of time.
- Now parallelize it
 - Set up cores
 - Only code change: use `parLapply()` instead of `lapply()`

Parallel Computing

- Set-up is a lot of work!
- For a large problem this could save a lot of time.
- Now parallelize it
 - Set up cores
 - Only code change: use `parLapply()` instead of `lapply()`

```
library(parallel)
cores <- detectCores()
cluster <- makeCluster(cores - 1)
results <- parLapply(cluster, X = c(250, 250, 250),
                     fun = parallel.function, data = iris)
temp.vector <- sapply(results, function(result) {result$tot.withinss})
result <- results[[which.min(temp.vector)]]
print(result)
```

Parallel Computing

```
## K-means clustering with 3 clusters of sizes 38, 62, 50
##
## Cluster means:
##   Sepal.Length Sepal.Width Petal.Length Petal.Width
## 1      6.850000      3.073684      5.742105      2.071053
## 2      5.901613      2.748387      4.393548      1.433871
## 3      5.006000      3.428000      1.462000      0.246000
##
## Clustering vector:
##   [1] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
##  [36] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 2 2 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
##  [71] 2 2 2 2 2 2 2 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 2 1 1 1
## [106] 1 2 1 1 1 1 1 1 2 2 1 1 1 1 2 1 2 1 2 1 1 2 2 1 1 1 1 1 2 1 1 1 2 1
## [141] 1 1 2 1 1 1 2 1 1 2
##
## Within cluster sum of squares by cluster:
## [1] 23.87947 39.82097 15.15100
## (between_SS / total_SS =  88.4 %)
##
## Available components:
```

80/84

Parallel Computing

- Compare computation time when parallelized

```
parTime <- microbenchmark({  
  library(parallel)  
  cores <- detectCores()  
  cluster <- makeCluster(cores - 1)  
  results <- parLapply(cluster, X = c(25000, 25000, 25000),  
                        fun = parallel.function, data = iris)  
  temp.vector <- sapply(results, function(result) {result$tot.withinss})  
  result <- results[[which.min(temp.vector)]]  
}, times = 10, unit = "s")  
  
straightTime <- microbenchmark({  
  clus <- kmeans(iris, centers = 3, nstart = 75000)  
}, times = 10, unit = "s")
```

Parallel Computing

parTime

```
## Unit: seconds
##
## {      library(parallel)      cores <- detectCores()      cluster <- makeCluster(cores - 1)
##      min      lq      mean      median      uq      max neval
##  2.479656 2.499244 2.513047 2.511559 2.529938 2.544348    10
```

straightTime

```
## Unit: seconds
##
##                                     expr      min
## {      clus <- kmeans(iris, centers = 3, nstart = 75000) } 3.489261
##      lq      mean      median      uq      max neval
##  3.502473 3.560645 3.562536 3.591032 3.647707    10
```

Recap!

- Parallel Computing can speed up computations
- A lot of up front work
- Can only use when process can be done separately
- Many other ways to speed up R as well (see Microsoft R)

What do we want to be able to do?

- Restructure Data/Clean Data
- Streamline repeated sections of code
- Improve efficiency of code
- Write custom functions to simplify code
- Thanks for coming!